

R Notebook

Part-A

This chunk help us to load the dataset.

```
myFile <- 'D:/ahmed al/shz.txt' # the local file path to my research prospect
us
# fill = TRUE b/c rows are of unequal length
#fill: Sometimes, we may get a file that contains the
# unequal length of rows, and we have to add blank spaces to that missing val
ues.
dat <- read.table(myFile, header = FALSE, fill = TRUE)
```

Loading our required packages.

```
library(dplyr) # for data wrangling

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

library(tidytext) # for Natural Language Processing

## Warning: package 'tidytext' was built under R version 4.1.2

library(stringr) # to deal with strings
library(wordcloud) # to for generating word clouds

## Warning: package 'wordcloud' was built under R version 4.1.2

## Loading required package: RColorBrewer

library(knitr) # for tables, It combines many features into one package

## Warning: package 'knitr' was built under R version 4.1.2

library(DT) # for dynamic tables

## Warning: package 'DT' was built under R version 4.1.2

library(tidyr)
```

```
## Warning: package 'tidyr' was built under R version 4.1.2
```

```
#Tools to help to create tidy data, where each column is a variable,  
#each row is an observation, and each cell contains a single value.
```

Gather help us to create a single coloumn of all the words.

```
tidy_dat <- tidyr::gather(dat, key, word) %>% select(word)  
head(tidy_dat)
```

```
##           word  
## 1         Zyaed  
## 2 Understanding  
## 3           Born  
## 4           Life,  
## 5       Through  
## 6           The
```

length tells us about the total number of words(also known as tokens in NLP)

```
tidy_dat$word %>% length() #there are 2832 tokens in my document
```

```
## [1] 2832
```

Unique help us to identify the unique words.

```
unique(tidy_dat$word) %>% length() # 695 words are unique
```

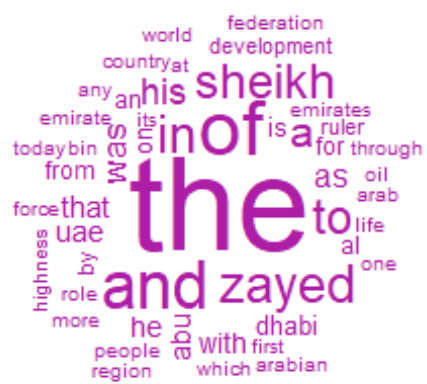
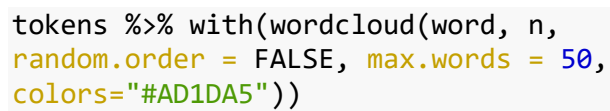
```
## [1] 695
```

Following commands help us to count the total number of words.

```
tokens <- tidy_dat %>%  
unnest_tokens(word, word) %>%  
dplyr::count(word, sort = TRUE) %>%  
ungroup()  
tokens %>% head(10)
```

```
##      word    n  
## 1    the 137  
## 2     of   73  
## 3    and   62  
## 4     in   45  
## 5   zayed  45  
## 6     to  40  
## 7 sheikh  39  
## 8      a   32  
## 9    his  25  
## 10   was  22
```

Word cloud is graph that shows the most frequent words by there size. Following commands use the total number of repetitions to create the word cloud.



It removes the stopping words like is,or,am etc from the text data.

```
# remove stop words
data("stop_words")
tokens_clean <- tokens %>%
  anti_join(stop_words, by = "word")
```

It removes the numbers from the text data.

```
# remove numbers
nums <- tokens_clean %>% filter(str_detect(word, "[0-9]")) %>%
  select(word) %>% unique()
tokens_clean <- tokens_clean %>%
  anti_join(nums, by = "word")
```

Following code removes the i.e and e.g from the data.

```
# remove other stop words
uni_sw <- data.frame(word = c("i.e", "e.g."))
tokens_clean <- tokens_clean %>%
  anti_join(uni_sw, by = "word")
```

Creating the word cloud again without the numbers and stop words.

```
# define a nice color palette
pal <- brewer.pal(8, "Dark2")
# plot the 50 most common words
tokens_clean %>%
  with(wordcloud(word, n, random.order = FALSE, max.words = 50,
    colors=pal))
```



Following

commands show the clean data with word frequencies.

```
tokens_clean %>%
  DT::datatable()
```

Show entries Search:

	word	n
1	zayed	45
2	sheikh	39
3	uae	17
4	abu	14
5	dhabi	12
6	al	10
7	development	8
8	ruler	7
9	country	6
10	life	6

Showing 1 to 10 of 491 entries Previous 1 2 3 4 5 ... 50 Next

```
head(tokens_clean)
```

```
##      word  n
## 1  zayed 45
## 2 sheikh 39
## 3   uae 17
## 4   abu 14
## 5 dhabi 12
## 6    al 10
```

R sentiment packages help us to identify the sentiments from the text data. We created text and checked the sentiments with Rsentiments packages.

```
# Load the Library
library(RSentiment)

## Warning: package 'RSentiment' was built under R version 4.1.3

# Module Name 69
calculate_total_presence_sentiment(c("This is a good text", "This is a bad te
xt",
"This is a really bad text", "This is horrible"))

## [1] "Processing sentence: this is a good text"
## [1] "Processing sentence: this is a bad text"
## [1] "Processing sentence: this is a really bad text"
## [1] "Processing sentence: this is horrible"

##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] "Sarcasm" "Negative" "Very Negative" "Neutral" "Positive" "Very Posit
ive"
## [2,] "0"      "3"      "0"      "0"      "1"      "0"
```

This is telling us about the positive and negative sentiments of the sentences.

```
calculate_sentiment(c("This is a good text",
  "This is a bad text",
  "This is a really bad text", "This is horrible"))

## [1] "Processing sentence: this is a good text"
## [1] "Processing sentence: this is a bad text"
## [1] "Processing sentence: this is a really bad text"
## [1] "Processing sentence: this is horrible"

##           text sentiment
## 1   This is a good text  Positive
## 2   This is a bad text   Negative
## 3 This is a really bad text Negative
## 4   This is horrible     Negative

calculate_score(c("This is a good text",
  "This is a bad text",
  "This is a really bad text",
  "This is horrible"))
```

```

## [1] "Processing sentence: this is a good text"
## [1] "Processing sentence: this is a bad text"
## [1] "Processing sentence: this is a really bad text"
## [1] "Processing sentence: this is horrible"

## [1] 1 -1 -1 -1

library(tidytext)
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.1.3

## -- Attaching packages ----- tidyverse 1.
3.1 --

## v ggplot2 3.3.5      v purrr 0.3.4
## v tibble 3.1.2      v forcats 0.5.1
## v readr 2.1.1

## Warning: package 'ggplot2' was built under R version 4.1.1
## Warning: package 'readr' was built under R version 4.1.2

## -- Conflicts ----- tidyverse_conflict
s() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()

library(janeaustenr)

## Warning: package 'janeaustenr' was built under R version 4.1.2

library(stringr)
library(wordcloud)
library(reshape2)

##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
## smiths

library(textdata)

## Warning: package 'textdata' was built under R version 4.1.2

```

There are dictionaries to identify the sentiments of the text data. There are three different libraries, *affin*, *bing* and *nrc*. *Affin* package gives sentiments in the numeric form and *bing* identify the positive and negative sentiments. and *nrc* can identify the different kind of the sentiments like trust, fear, sadness, anger, etc..

```
get_sentiments("afinn")
```



```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon      -2
## 2 abandoned    -2
## 3 abandons     -2
## 4 abducted     -2
## 5 abduction    -2
## 6 abductions   -2
## 7 abhor        -3
## 8 abhorred     -3
## 9 abhorrent    -3
## 10 abhors      -3
## # ... with 2,467 more rows
```

```
get_sentiments("bing")
```

```
## # A tibble: 6,786 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 2-faces   negative
## 2 abnormal negative
## 3 abolish  negative
## 4 abominable negative
## 5 abominably negative
## 6 abominate negative
## 7 abomination negative
## 8 abort     negative
## 9 aborted   negative
## 10 aborts    negative
## # ... with 6,776 more rows
```

```
get_sentiments("nrc")
```

```
## # A tibble: 13,875 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus    trust
## 2 abandon    fear
## 3 abandon    negative
## 4 abandon    sadness
## 5 abandoned  anger
## 6 abandoned  fear
## 7 abandoned  negative
## 8 abandoned  sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,865 more rows
```

```
data(sentiments)
#dataset structure
str(sentiments)

## tibble [6,786 x 2] (S3: tbl_df/tbl/data.frame)
## $ word      : chr [1:6786] "2-faces" "abnormal" "abolish" "abominable" ...
## $ sentiment: chr [1:6786] "negative" "negative" "negative" "negative" ...
```

We are saving the affins as the `affin_lexicon` variable. and we are doing same thing for the `afinn`, `bing`, `nrc`.

```
afinn_lexicon <- get_sentiments("afinn")
head(afinn_lexicon)
```

```
## # A tibble: 6 x 2
##   word      value
##   <chr>     <dbl>
## 1 abandon    -2
## 2 abandoned  -2
## 3 abandons   -2
## 4 abducted   -2
## 5 abduction  -2
## 6 abductions -2
```

#NRC

```
nrc_lexicon <- get_sentiments("nrc")
head(nrc_lexicon)
```

```
## # A tibble: 6 x 2
##   word      sentiment
##   <chr>     <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
```

#BING

```
bing_lexicon <- get_sentiments("bing")
head(bing_lexicon)
```

```
## # A tibble: 6 x 2
##   word      sentiment
##   <chr>     <chr>
## 1 2-faces   negative
## 2 abnormal  negative
## 3 abolish   negative
## 4 abominable negative
## 5 abominably negative
## 6 abominate  negative
```

We are using Jane Austens data of books for the further analysis.

```
tidy_books <- austen_books() %>%
  filter(book == "Emma") %>%
  group_by(book) %>%
  ## Using row_number() with mutate() will create a column of consecutive numbers. The row_number() function
  #is useful for creating an identification number (an ID variable). It is also useful for labeling each observation by a
  #grouping variable.
  ##cumsum() function in R Language is used to calculate the cumulative sum of the vector passed as
  ## argument
  # str_detect function returns a logical value (i.e. FALSE or TRUE),
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text, regex("^chapter [\\divxlc]",
         ignore_case = TRUE)))) %>% ungroup() %>%
  unnest_tokens(word, text)
```

following codes separate the words with joy only in nrc library.

```
nrc_joy <- get_sentiments("nrc") %>%
  filter(sentiment == "joy")
#Summarize the usage of `joy` words
tidy_books %>%
  semi_join(nrc_joy) %>%
  count(word, sort = T)

## Joining, by = "word"

## # A tibble: 301 x 2
##   word      n
##   <chr>   <int>
## 1 good    359
## 2 friend  166
## 3 hope    143
## 4 happy   125
## 5 love    117
## 6 deal     92
## 7 found    92
## 8 present  89
## 9 kind     82
## 10 happiness 76
## # ... with 291 more rows

tidy_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text, regex("^chapter [\\divxlc]",
```

```
ignore_case = TRUE)))) %>%

ungroup() %>%
unnest_tokens(word, text)
```

In the following codes we are generating the random data and using it for the sentimental analysis.

```
## Creating identification number to represent 50 individual people
ID <- c(1:20)
## Creating sex variable (10 males/10 females)
Sex <- rep(c("male", "female"), 10) # rep stands for replicate
## Creating age variable (20-39 year olds)
Age <- c(26, 25, 39, 37, 31, 34, 34, 30, 26, 33,
  39, 28, 26, 29, 33, 22, 35, 23, 26, 36)
## Creating a dependent variable called Score
Score <- c(0.010, 0.418, 0.014, 0.090, 0.061, 0.328, 0.656, 0.002, 0.639, 0.1
73,
  0.076, 0.152, 0.467, 0.186, 0.520, 0.493, 0.388, 0.501, 0.800, 0.482)
## Creating a unified dataset that puts together all variables
## tibble is a simple dataframe
data <- tibble(ID, Sex, Age, Score)
## group by sex
data %>%
  group_by(Sex) %>%
  summarize(m = mean(Score), # calculates the mean
    s = sd(Score), # calculates the standard deviation
    n = n()) %>% # calculates the total number of observations
  ungroup()

## # A tibble: 2 x 4
##   Sex      m      s      n
##   <chr> <dbl> <dbl> <int>
## 1 female 0.282 0.184    10
## 2 male   0.363 0.300    10

## `summarise()` ungrouping output (override with `.groups` argument)
# A tibble: 2 x 4
##x m s n
##hr> <dbl> <dbl> <int>
##1 female 0.282 0.184 10
##2 male 0.363 0.300 10
## mutate() and group_by()
data %>%
  group_by(Sex) %>%
  mutate(m = mean(Score)) %>% # calculates mean score by Sex
  ungroup()

## # A tibble: 20 x 5
##       ID Sex      Age Score      m
##   <int> <chr> <dbl> <dbl> <dbl>
```

```
## 1      1 male      26 0.01  0.363
## 2      2 female    25 0.418 0.282
## 3      3 male      39 0.014 0.363
## 4      4 female    37 0.09  0.282
## 5      5 male      31 0.061 0.363
## 6      6 female    34 0.328 0.282
## 7      7 male      34 0.656 0.363
## 8      8 female    30 0.002 0.282
## 9      9 male      26 0.639 0.363
## 10     10 female   33 0.173 0.282
## 11     11 male     39 0.076 0.363
## 12     12 female   28 0.152 0.282
## 13     13 male     26 0.467 0.363
## 14     14 female   29 0.186 0.282
## 15     15 male     33 0.52  0.363
## 16     16 female   22 0.493 0.282
## 17     17 male     35 0.388 0.363
## 18     18 female   23 0.501 0.282
## 19     19 male     26 0.8    0.363
## 20     20 female   36 0.482 0.282
```

```
janeaustensentiment <- tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(book, index = linenumber %/% 100, sentiment) %>%

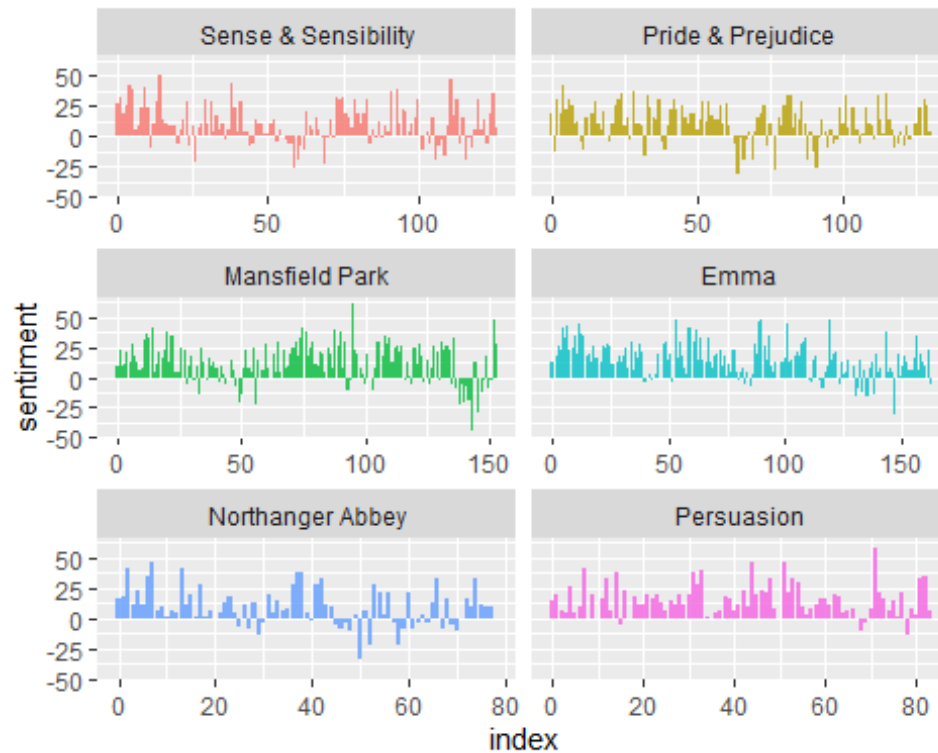
  spread(sentiment, n, fill = 0) %>%

  mutate(sentiment = positive - negative)

## Joining, by = "word"
```

Following plot shows the data for the sentiments for different chapters of the Jane book.

```
ggplot(data = janeaustensentiment, mapping = aes(x = index, y =
sentiment, fill = book)) +
  geom_bar(alpha = 0.8, stat = "identity", show.legend = FALSE) +
  facet_wrap(facets = ~ book, ncol = 2, scales = "free_x")
```



Part-B

I am using texts available on this website.

<https://www.toppr.com/guides/essays/freedom-essay/> . Lets load the dataset.

```
myFile <- 'D:/ahmed al/txt.txt'
dat <- read.table(myFile, header = FALSE, fill = TRUE)
head(dat)
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10
## 1	SIGNUP									
## 2	LOGIN									
## 3	Search	for	a	topic						
## 4	English	>	Essays	>	Freedom	Essay	for	Students	and	Children
## 5	Essays									
## 6	Freedom	Essay	for	Students	and	Children				

let's reshape the dataset for the sentimental analysis.

```
tidy_dat <- tidyr::gather(dat, key, word) %>% select(word)
head(tidy_dat)
```

	word
## 1	SIGNUP
## 2	LOGIN

```
## 3 Search
## 4 English
## 5 Essays
## 6 Freedom
```

Counting the words in the text file.

```
tokens <- tidy_dat %>%
unnest_tokens(word, word) %>%
dplyr::count(word, sort = TRUE) %>%
ungroup()
tokens %>% head(10)
```

```
##      word  n
## 1 freedom 37
## 2    the  35
## 3     of  34
## 4    and  28
## 5    is  19
## 6  essay  17
## 7  right  17
## 8     to  17
## 9    for  16
## 10    it  14
```

Let's remove the stop words from the dataset.

```
# remove stop words
data("stop_words")
tokens_clean <- tokens %>%
  anti_join(stop_words, by = "word")
```

Let's remove the numbers from the dataset.

```
# remove numbers
nums <- tokens_clean %>% filter(str_detect(word, "[0-9]")) %>%
  select(word) %>% unique()
tokens_clean <- tokens_clean %>%
  anti_join(nums, by = "word")
```

Let's create the word cloud of the data.

```
pal <- brewer.pal(8, "Dark2")
# plot the 50 most common words
tokens_clean %>%
  with(wordcloud(word, n, random.order = FALSE, max.words = 50,
    colors=pal))
```



Part-C

```
#install.packages("SentimentAnalysis")
library(SentimentAnalysis)

## Warning: package 'SentimentAnalysis' was built under R version 4.1.2
##
## Attaching package: 'SentimentAnalysis'
## The following object is masked from 'package:base':
##
##      write

sentiment <- analyzeSentiment("Yeah, this was a great soccer game for the Ger
man team!")
sentiment$PositivityGI

## [1] 0.3333333

#In case you had any errors, run this line of code:
#install.packages("SnowballC")
library(SnowballC)

## Warning: package 'SnowballC' was built under R version 4.1.1

sentiment$NegativityGI
```



```
## [1] 0
sentiment$WordCount
## [1] 6
```

Positive score is 0.3333 and negative score is zero. The word count is 6.

```
documents <- c ("Wow, I really like the new light sabers!",
                "That book was excellent.",
                "R is a fantastic language.",
                "The service in this restaurant was miserable.",
                "This is neither positive or negative.",
                "The waiter forget about my dessert -- what poor service!")
sentiment <- analyzeSentiment(documents)
sentiment$SentimentQDAP
## [1] 0.3333333 0.5000000 0.5000000 -0.3333333 0.0000000 -0.4000000
convertToDirection(sentiment$SentimentQDAP)
## [1] positive positive positive negative neutral negative
## Levels: negative neutral positive
convertToBinaryResponse(sentiment$SentimentQDAP)
## [1] positive positive positive negative positive negative
## Levels: negative positive
```

This code gets the sentiment score for the sentences and convert the sentiment scores into the binary responses, and converts it into factors.

```
response <- c(+1, +1, +1, -1, 0, -1)
compareToResponse(sentiment, response)
## Warning in cor(sentiment, response): the standard deviation is zero
## Warning in cor(x, y): the standard deviation is zero
## Warning in cor(x, y): the standard deviation is zero
## Warning in cor(sentiment, response): the standard deviation is zero
##
##          WordCount  SentimentGI  NegativityGI  PositivityGI
## cor          -0.18569534  0.990011498 -9.974890e-01  0.94295
## 4167
## cor.t.statistic -0.37796447 14.044046450 -2.816913e+01  5.66470
## 5543
## cor.p.value     0.72465864  0.000149157  9.449687e-06  0.00478
## 8521
## lm.t.value     -0.37796447 14.044046450 -2.816913e+01  5.66470
## 5543
```

## r.squared 2562	0.03448276	0.980122766	9.949843e-01	0.88916
## RMSE 4032	3.82970843	0.450102869	1.186654e+00	0.71362
## MAE 6667	3.33333333	0.400000000	1.100000e+00	0.66666
## Accuracy 6667	0.66666667	1.000000000	6.666667e-01	0.66666
## Precision NaN	NaN	1.000000000	NaN	
## Sensitivity 0000	0.00000000	1.000000000	0.000000e+00	0.00000
## Specificity 0000	1.00000000	1.000000000	1.000000e+00	1.00000
## F1 NaN	NaN	1.000000000	NaN	
## BalancedAccuracy 0000	0.50000000	1.000000000	5.000000e-01	0.50000
## avg.sentiment.pos.response 6667	3.25000000	0.333333333	8.333333e-02	0.41666
## avg.sentiment.neg.response 0000	4.00000000	-0.633333333	6.333333e-01	0.00000
## LM	SentimentHE	NegativityHE	PositivityHE	Sentiment
## cor 55	0.4152274	-0.083045480	0.3315938	0.73704
## cor.t.statistic 42	0.9128709	-0.166666667	0.7029595	2.18111
## cor.p.value 66	0.4129544	0.875718144	0.5208394	0.09462
## lm.t.value 42	0.9128709	-0.166666667	0.7029595	2.18111
## r.squared 61	0.1724138	0.006896552	0.1099545	0.54323
## RMSE 78	0.8416254	0.922958207	0.8525561	0.72341
## MAE 33	0.7500000	0.888888889	0.8055556	0.63333
## Accuracy 33	0.6666667	0.666666667	0.6666667	0.83333
## Precision 00	NaN	NaN	NaN	1.00000
## Sensitivity 00	0.0000000	0.000000000	0.0000000	0.50000
## Specificity 00	1.0000000	1.000000000	1.0000000	1.00000
## F1 67	NaN	NaN	NaN	0.66666
## BalancedAccuracy 00	0.5000000	0.500000000	0.5000000	0.75000

```
## avg.sentiment.pos.response 0.1250000 0.08333333 0.2083333 0.2500000
## avg.sentiment.neg.response 0.0000000 0.00000000 0.0000000 -0.1000000
##
## NegativityLM PositivityLM RatioUncertaintyLM
## cor -0.40804713 0.6305283 NA
## cor.t.statistic -0.89389841 1.6247248 NA
## cor.p.value 0.42189973 0.1795458 NA
## lm.t.value -0.89389841 1.6247248 NA
## r.squared 0.16650246 0.3975659 NA
## RMSE 0.96186547 0.7757911 0.9128709
## MAE 0.92222222 0.7222222 0.8333333
## Accuracy 0.66666667 0.6666667 0.6666667
## Precision NaN NaN NaN
## Sensitivity 0.00000000 0.0000000 0.0000000
## Specificity 1.00000000 1.0000000 1.0000000
## F1 NaN NaN NaN
## BalancedAccuracy 0.50000000 0.5000000 0.5000000
## avg.sentiment.pos.response 0.08333333 0.3333333 0.0000000
## avg.sentiment.neg.response 0.10000000 0.0000000 0.0000000
##
## SentimentQDAP NegativityQDAP PositivityQDAP
## cor 0.9865356369 -0.944339551 0.942954167
## cor.t.statistic 12.0642877257 -5.741148345 5.664705543
## cor.p.value 0.0002707131 0.004560908 0.004788521
## lm.t.value 12.0642877257 -5.741148345 5.664705543
## r.squared 0.9732525629 0.891777188 0.889162562
## RMSE 0.5398902495 1.068401367 0.713624032
## MAE 0.4888888889 1.011111111 0.666666667
## Accuracy 1.0000000000 0.666666667 0.666666667
## Precision 1.0000000000 NaN NaN
## Sensitivity 1.0000000000 0.000000000 0.000000000
## Specificity 1.0000000000 1.000000000 1.000000000
## F1 1.0000000000 NaN NaN
## BalancedAccuracy 1.0000000000 0.500000000 0.500000000
## avg.sentiment.pos.response 0.3333333333 0.083333333 0.416666667
## avg.sentiment.neg.response -0.3666666667 0.366666667 0.000000000
```

We created the exact responses for the sentences and compared them with the sentiment score by our analysis and got all the scores by comparison.

Part-D

```
##1.2 The unnest_tokens function
#Emily Dickinson wrote some lovely text in her time.
text <- c("Because I could not stop for Death -",
          "He kindly stopped for me -",
          "The Carriage held but just Ourselves -",
```

```

      "and Immortality")
text

## [1] "Because I could not stop for Death -"
## [2] "He kindly stopped for me -"
## [3] "The Carriage held but just Ourselves -"
## [4] "and Immortality"

##This is a typical character vector that we might want to analyze.
##To turn it into a tidy text dataset,
##we first need to put it into a data frame.
library(dplyr)
text_df <- tibble(line = 1:4, text = text)
text_df

## # A tibble: 4 x 2
##   line text
##   <int> <chr>
## 1     1 Because I could not stop for Death -
## 2     2 He kindly stopped for me -
## 3     3 The Carriage held but just Ourselves -
## 4     4 and Immortality

##Within our tidy text framework, we need to both break the text into
##individual tokens (a process called tokenization) and transform it to
##a tidy data structure. To do this, we use tidytext's unnest_tokens() functio
n.
library(tidytext)
text_df %>%
  unnest_tokens(word, text)

## # A tibble: 20 x 2
##   line word
##   <int> <chr>
## 1     1 because
## 2     1 i
## 3     1 could
## 4     1 not
## 5     1 stop
## 6     1 for
## 7     1 death
## 8     2 he
## 9     2 kindly
## 10    2 stopped
## 11    2 for
## 12    2 me
## 13    3 the
## 14    3 carriage
## 15    3 held
## 16    3 but
## 17    3 just

```

```
## 18      3 ourselves
## 19      4 and
## 20      4 immortality
```

Just appear at the line 3.

##1.3 Tidying the works of Jane Austen published novels

```
library(janeaustenr)
library(dplyr)
library(stringr)
original_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text,
                                     regex("^chapter [\\divxlc]",
                                           ignore_case = TRUE)))) %>%
  ungroup()
original_books

## # A tibble: 73,422 x 4
##   text                book          linenumber chapter
##   <chr>              <fct>          <int>    <int>
## 1 "SENSE AND SENSIBILITY" Sense & Sensibility      1         0
## 2 ""                Sense & Sensibility      2         0
## 3 "by Jane Austen"   Sense & Sensibility      3         0
## 4 ""                Sense & Sensibility      4         0
## 5 "(1811)"           Sense & Sensibility      5         0
## 6 ""                Sense & Sensibility      6         0
## 7 ""                Sense & Sensibility      7         0
## 8 ""                Sense & Sensibility      8         0
## 9 ""                Sense & Sensibility      9         0
## 10 "CHAPTER 1"       Sense & Sensibility     10         1
## # ... with 73,412 more rows
```

CHAPTER 1 appears at the line 10 and chapter 1.

#To work with this as a tidy dataset, we need to restructure it in the one-to-ken-per-row format,

#which as we saw earlier is done with the unnest_tokens() function.

```
library(tidytext)
tidy_books <- original_books %>%
  unnest_tokens(word, text)
tidy_books

## # A tibble: 725,055 x 4
##   book          linenumber chapter word
##   <fct>          <int>    <int> <chr>
## 1 Sense & Sensibility      1         0 sense
## 2 Sense & Sensibility      1         0 and
## 3 Sense & Sensibility      1         0 sensibility
## 4 Sense & Sensibility      3         0 by
```

```
## 5 Sense & Sensibility      3      0 jane
## 6 Sense & Sensibility      3      0 austen
## 7 Sense & Sensibility      5      0 1811
## 8 Sense & Sensibility     10      1 chapter
## 9 Sense & Sensibility     10      1 1
## 10 Sense & Sensibility     13      1 the
## # ... with 725,045 more rows
```

#This function uses the tokenizers package to separate each line of text in the original data frame into tokens.

remove stop words (kept in the tidytext dataset stop_words)

```
data(stop_words)
tidy_books <- tidy_books %>%
  anti_join(stop_words)
```

```
## Joining, by = "word"
```

Tokenization means splitting the sentences into words and counting the appearance of the words. It remove all the stop words from the data.

##use dplyr's count() to find the most common words in all the books as a whole

```
tidy_books %>%
  count(word, sort = TRUE)
```

```
## # A tibble: 13,914 x 2
##   word      n
##   <chr> <int>
## 1 miss   1855
## 2 time   1337
## 3 fanny   862
## 4 dear    822
## 5 lady    817
## 6 sir     806
## 7 day     797
## 8 emma    787
## 9 sister  727
## 10 house  699
## # ... with 13,904 more rows
```


These are the top 5 common words. miss time fanny dear and lady.

##create a visualization of the most common words

```
library(ggplot2)
```

```
tidy_books %>%
  count(word, sort = TRUE) %>%
  filter(n > 300) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word)) +
```

```
geom_col() +  
labs(y = NULL)
```



The above code counts the words frequency, and filter the words that occur more than 300 times. After that we reorder the data and plot them on the y axis in the graph. The word 'miss' is the most occurring and 'eyes' is the least occurring.